

Game Playing

Today's class

- Game playing
 - State of the art and resources
 - Framework
- Game trees
 - Minimax
 - Alpha-beta pruning
 - Adding randomness

Why study games?

- Clear criteria for success
- Offer an opportunity to study problems involving {hostile, adversarial, competing} agents.
- Historical reasons
- Fun
- Interesting, hard problems which require minimal “initial structure”
- Games often define very large search spaces
 - chess 35^{100} nodes in search tree, 10^{40} legal states

State of the art

- How good are computer game players?
 - **Chess:**
 - Deep Blue beat Gary Kasparov in 1997
 - Garry Kasparov vs. Deep Junior (Feb 2003): tie!
 - Kasparov vs. X3D Fritz (November 2003): tie!
<http://www.thechessdrum.net/tournaments/Kasparov-X3DFritz/index.html>
 - Deep Fritz beat world champion Vladimir Kramnik (2006)
 - **Checkers:** Chinook (an AI program with a *very large* endgame database) is the world champion and can provably never be beaten. Retired in 1995
 - **Go:** Computer players have finally reached tournament-level play
 - **Bridge:** “Expert-level” computer players exist (but no world champions yet!)
- Good places to learn more:
 - <http://www.cs.ualberta.ca/~games/>
 - <http://www.cs.unimass.nl/icga>

Typical case

- 2-person game
- Players alternate moves
- **Zero-sum**: one player's loss is the other's gain
- **Perfect information**: both players have access to complete information about the state of the game. No information is hidden from either player.
- No chance (e.g., using dice) involved
- Examples: Tic-Tac-Toe, Checkers, Chess, Go, Nim, Othello
- Not: Bridge, Solitaire, Backgammon, ...

How to play a game

- A way to play such a game is to:
 - Consider all the legal moves you can make
 - Compute the new position resulting from each move
 - Evaluate each resulting position and determine which is best
 - Make that move
 - Wait for your opponent to move and repeat
- Key problems are:
 - Representing the “board”
 - Generating all legal next boards
 - Evaluating a position

Evaluation function

- **Evaluation function** or **static evaluator** is used to **evaluate** the “goodness” of a game position.
 - Contrast with **heuristic search** where the evaluation function was a non-**negative estimate** of the cost from the start node to a goal and passing through the given node
- The zero-sum assumption allows us to use a single evaluation function to describe the goodness of a board with respect to both players.
 - $f(n) \gg 0$: position n good for me and bad for you
 - $f(n) \ll 0$: position n bad for me and good for you
 - $f(n)$ near 0 : position n is a **neutral position**
 - $f(n) = +\text{infinity}$: win for me
 - $f(n) = -\text{infinity}$: win for you


Evaluation function examples

- Example of an evaluation function for Tic-Tac-Toe:
$$f(n) = [\text{\# of 3-lengths open for me}] - [\text{\# of 3-lengths open for you}]$$

where a 3-length is a complete row, column, or diagonal
- Alan Turing's function for chess
 - $f(n) = w(n)/b(n)$ where $w(n)$ = sum of the point value of white's pieces and $b(n)$ = sum of black's
- Most evaluation functions are specified as a weighted sum of position features:
$$f(n) = w_1 * \text{feat}_1(n) + w_2 * \text{feat}_2(n) + \dots + w_n * \text{feat}_k(n)$$
- Example features for chess are piece count, piece placement, squares controlled, etc.
- Deep Blue had over 8000 features in its evaluation function


Game trees

- Problem spaces for typical games are represented as trees
- Root node represents the current board configuration; player must decide the best single move to make next
- **Static evaluator function** rates a board position. $f(\text{board}) = \text{real number}$ with $f > 0$ “white” (me), $f < 0$ for black (you)
- Arcs represent the possible legal moves for a player
- If it is **my turn** to move, then the root is labeled a “**MAX**” node; otherwise it is labeled a “**MIN**” node, indicating **my opponent's turn**.
- Each level of the tree has nodes that are all MAX or all MIN; nodes at level i are of the opposite kind from those at level $i+1$



MinMax - Overview

- Search tree
 - *Squares* represent decision states (ie- after a move)
 - *Branches* are decisions (ie- the move)
 - Start at root
 - Nodes at end are leaf nodes
 - Ex: Tic-Tac-Toe (symmetrical positions removed)




- Unlike binary trees can have any number of children
 - Depends on the game situation
- Levels usually called *plies* (a *ply* is one level)
 - Each ply is where "turn" switches to other player
- Players called *Min* and *Max* (next)

Minimax procedure

- Create **start node as a MAX node** with current board configuration
- **Expand nodes down to some depth** (a.k.a. **ply**) of lookahead in the game
- **Apply the evaluation function at each of the leaf nodes**
- “Back up” values for each of the non-leaf nodes until a value is computed for the root node
 - At **MIN nodes**, the backed-up value is **the minimum of the values** associated with its children.
 - At **MAX nodes**, the backed-up value is the **maximum of the values** associated with its children.
- Pick the operator associated with the child node whose backed-up value determined the value at the root


Minimax

R




Minimax

R





Minimax

R




Minimax Algorithm





Minimax


Picking my best move against your best move



$minimax(s) =$


$$\begin{cases} \text{utility}(s) & \text{if } \text{terminal}(s) \\ \max_{a \in \text{action}(s)} minimax(\text{result}(s, a)) & \text{if } \text{player}(s) = \text{MAX} \\ \min_{a \in \text{action}(s)} minimax(\text{result}(s, a)) & \text{if } \text{player}(s) = \text{MIN} \end{cases}$$

Partial Game Tree for Tic-Tac-Toe



- $f(n) = +1$ if the position is a win for X.
- $f(n) = -1$ if the position is a win for O.
- $f(n) = 0$ if the position is a draw.

Minimax Tree




Minimax Discussion

- ▶ Complete depth first exploration
- ▶ Depth m with b legal moves. $O(b^m)$
- ▶ Space complexity (memory) $O(bm)$
- ▶ Chess: $m \approx 35$; on average: $50 \leq b \leq 100$
- ▶ Impractical for most games, but basis of other algs.

Alpha-beta pruning

- We can improve on the performance of the minimax algorithm through **alpha-beta pruning**
- Basic idea: *“If you have an idea that is surely bad, don't take the time to see how truly awful it is.”* -- Pat Winston




- We don't need to compute the value at this node.
- No matter what it is, it can't affect the value of the root node.

Alpha-beta pruning


- Traverse the search tree in depth-first order
- At each **MAX** node n , **alpha(n)** = maximum value found so far
- At each **MIN** node n , **beta(n)** = minimum value found so far
 - Note: The alpha values start at -infinity and only increase, while beta values start at +infinity and only decrease.
- **Beta cutoff:** Given a MAX node n , cut off the search below n (i.e., don't generate or examine any more of n 's children) if $\text{alpha}(n) \geq \text{beta}(i)$ for some MIN node ancestor i of n .
- **Alpha cutoff:** stop searching below MIN node n if $\text{beta}(n) \leq \text{alpha}(i)$ for some MAX node ancestor i of n .

Alpha-beta example



Alpha-Beta pruning

Intuition



Do we need to expand all nodes?

$$\begin{aligned} \text{minimax}(\text{root}) &= \max(\underline{\min(3, 12, 8)}, \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \\ &= 3 \end{aligned}$$

Do we need z ?

Alpha-beta algorithm

```
function MAX-VALUE (state,  $\alpha$ ,  $\beta$ )
    ;;  $\alpha$  = best MAX so far;  $\beta$  = best MIN
    if TERMINAL-TEST (state) then return UTILITY(state)
    v :=  $-\infty$ 
    for each s in SUCCESSORS (state) do
        v := MAX (v, MIN-VALUE (s,  $\alpha$ ,  $\beta$ ))
        if v  $\geq$   $\beta$  then return v
         $\alpha$  := MAX ( $\alpha$ , v)
    end
    return v
```

```
function MIN-VALUE (state,  $\alpha$ ,  $\beta$ )
    if TERMINAL-TEST (state) then return UTILITY(state)
    v :=  $\infty$ 
    for each s in SUCCESSORS (state) do
        v := MIN (v, MAX-VALUE (s,  $\alpha$ ,  $\beta$ ))
        if v  $\leq$   $\alpha$  then return v
         $\beta$  := MIN ( $\beta$ , v)
    end
    return v
```

Two values:

- ▶ α = value of best choice so far for MAX (highest-value)
- ▶ β = value of best choice so far for MIN (lowest-value)
- ▶ Each node keeps track of its $[\alpha, \beta]$ values

Alpha-Beta Pruning Properties


- ▶ Pruning does not affect final outcome
- ▶ Sorting moves by result improves $\alpha - \beta$ performance
- ▶ Perfect ordering: $O(b^{\frac{m}{2}})$
- ▶ An exercise on **metareasoning**


```
def minimax(position, depth, maximizingPlayer):  
    if depth == 0 or game over in position:  
        return static evaluation of position
```


```
    if maximizingPlayer:  
        maxEval = -infinity  
        for child of position:  
            eval = minimax(child, depth - 1, false)  
            maxEval = max(maxEval, eval)  
        return maxEval
```


```
    else: # minimizingPlayer  
        minEval = +infinity  
        for child of position:  
            eval = minimax(child, depth - 1, true)  
            minEval = min(minEval, eval)  
        return minEval
```

```
return minimax(startPosition, 3, true)
```









```

k(position, depth, alpha, beta, maximizingPlayer)
    or game over in position
    static evaluation of position

```

```

Player
    -infinity
    child of position
    minimax(child, depth - 1, alpha, beta, false)
    = max(maxEval, eval)
    max(alpha, eval)
    <= alpha

```

Eval

```

+infinity
    child of position
    minimax(child, depth - 1, alpha, beta, true)
    = min(minEval, eval)
    min(beta, eval)
    <= alpha


```

Eval

```

11
return Position, 3, -∞, +∞, true)

```



Effectiveness of alpha-beta

- Alpha-beta is guaranteed to compute the same value for the root node as computed by minimax, with less or equal computation
- **Worst case:** no pruning, examining b^d leaf nodes, where each node has b children and a d -ply search is performed
- **Best case:** examine only $(2b)^{d/2}$ leaf nodes.
 - Result is you can search twice as deep as minimax!
- **Best case** is when each player's best move is the first alternative generated
- In Deep Blue, they found empirically that alpha-beta pruning meant that the average branching factor at each node was about 6 instead of about 35!

Games of chance

- Backgammon is a two-player game with **uncertainty**.
- Players roll dice to determine what moves to make.
- White has just rolled *5 and 6* and has four legal moves:
 - 5-10, 5-11
 - 5-11, 19-24
 - 5-10, 10-16
 - 5-11, 11-16
- Such games are good for exploring decision making in adversarial problems involving skill and luck.

